

# Parallelism and concurrency in Haskell

David Luposchinsky

2015-10-23

# Introduction

# Concurrency vs. parallelism

## Parallelism

- ▶ Goal: speedup
- ▶ Orthogonal to correctness
- ▶ Comparatively simple

## Concurrency

- ▶ Affects logical structure of a program
- ▶ Heavily interleaved with correctness
- ▶ Often very complicated

# Talk outline

1. **Parallelism**

Excellent case for laziness

2. **Concurrency with forks and mutable state**

Nice abstractions, language agnostic

3. **Transactional memory**

Unique to Haskell, envied by everyone else

# Parallelism

“Go faster”

## Checking primes

Problem: find primes in a list (e.g. for crypto)

```
isPrime :: Natural -> Bool
```

```
filterPrime :: [Natural] -> [Natural]
```

```
filterPrime = filter isPrime
```

```
primesBetween lo hi = filterPrime [lo..hi]
```

## Checking primes

Problem: find primes in a list (e.g. for crypto)

```
isPrime :: Natural -> Bool
```

```
filterPrime :: [Natural] -> [Natural]
```

```
filterPrime = filter isPrime
```

```
primesBetween lo hi = filterPrime [lo..hi]
```

Unfortunately, this won't parallelize well.



`map` is simple enough to parallelize (why?). Can we use that fact?

map is simple enough to parallelize (why?). Can we use that fact?

```
isPrime :: Natural -> Bool
```

```
filterPrime :: [Natural] -> [Natural]
```

```
filterPrime = map snd . filter fst . map (\n -> (isPrime n, n))  
    -- `isPrime` isolated in `map` argument  
    -- (using the decorate-operation-undecorate idiom)
```

```
primesBetween lo hi = filterPrime [lo..hi]
```

## What's in the list?

```
map (\n -> (isPrime n, n))  
    (take 1e5 [1e10..])
```

## What's in the list?

```
map (\n -> (isPrime n, n))  
    (take 1e5 [1e10..])
```

Thanks! Literally function applications of the lambda.

```
map (\n -> (isPrime n, n)) [1,2,3] = [ (\n -> (isPrime n, n)) 1  
                                         , (\n -> (isPrime n, n)) 2  
                                         , (\n -> (isPrime n, n)) 3 ]
```

# Forcing

- ▶ = Evaluation to a certain degree (e.g. *WHNF*)
- ▶ Primitive task that allows parallelism
- ▶ “I’ll probably need this, so maybe evaluate it”

# Evaluation strategies

```
-- Control.Parallel.Strategies
```

```
-- A way to evaluate 'a's
```

```
type Strategy a = a -> Eval a
```

```
withStrategy :: Strategy a -> a -> a
```

```
-- Primitives
```

```
r0    :: Strategy a -- "NOOP"
```

```
rseq  :: Strategy a -- Sequential forcing
```

```
rpar  :: Strategy a -- Parallel forcing
```

```
-- Combining functions
```

```
evalList      :: Strategy a -> Strategy [a]
```

```
evalTraversable :: Traversable t => Strategy a -> Strategy (t a)
```

```
evalTuple2    :: Strategy a -> Strategy b -> Strategy (a,b)
```

## Back to primes

```
filterPrime =  
  map snd . filter fst . map (\n -> (isPrime n, n))  
  --  
  -- List of thunks we'd like to force in parallel
```

## Back to primes

```
filterPrime =  
    map snd . filter fst . map (\n -> (isPrime n, n))  
    --  
    -- List of thunks we'd like to force in parallel  
  
filterPrimeParallel =  
    map snd . filter fst . parallelize . map (\n -> (isPrime n, n))  
    where  
        parallelize :: [(a, b)] -> [(a, b)]  
        parallelize = withStrategy tupleFstList  
  
        tupleFstList :: Strategy [(a, b)]  
        tupleFstList = parList parFst  
  
        parFst :: Strategy (a,b)  
        parFst = parTuple2 rseq r0
```



## Back to primes

```
filterPrime =
  map snd . filter fst . map (\n -> (isPrime n, n))
  --
  -- List of thunks we'd like to force in parallel

filterPrimeParallel =
  map snd . filter fst . parallelize . map (\n -> (isPrime n, n))
  where
    parallelize :: [(a, b)] -> [(a, b)]
    parallelize = withStrategy tupleFstList

    tupleFstList :: Strategy [(a, b)]
    tupleFstList = parList parFst

    parFst :: Strategy (a, b)
    parFst = parTuple2 rseq r0

filterPrimeParallel =
  map snd . filter fst . parallelize . map (\n -> (isPrime n, n))
  where
    parallelize = withStrategy (parList (parTuple2 rseq r0))
```

## Demo

```
. ./run  
[1 of 1] Compiling Main ( Primes.hs, Primes.o )  
Linking Primes ...  
Running sequentially  
45.38 s  
Running in parallel  
16.37 s
```

# Recap

`Control.Parallel` for easy parallelism:

1. Define strategy according to your needs
2. Apply strategy, done

# Threads are cheap

- ▶ Couple of hundred (heap) bytes per thread
- ▶ Literally millions are possible
- ▶ Don't be afraid of using them!

```
. ./run  
Creating thread ring of size 3000000  
Passing token 3000000 times  
1  
Time taken: 11.46 s
```

## Other parallelism libraries

- ▶ **monad-par**: Explicit dataflow networks for complicated (pure) computations
- ▶ **speculation**: the one-liner to solve all pure speculative parallelism needs

# Concurrency

# Basic primitives

```
forkIO :: IO () -> IO ThreadId

-- Inter-thread communication
data MVar a
newMVar  :: a -> IO (MVar a)
takeMVar :: MVar a -> IO a
putMVar  :: a -> MVar a -> IO ()
```

# MVar

- ▶ “Mutable variable”
- ▶ “Lock with data attached to it”
- ▶ Fully lazy, fully polymorphic
- ▶ Either full or empty
- ▶ Potentially blocking
  - ▶ write to full MVar
  - ▶ reading from empty MVar



## Example: network service

```
main :: IO ()
main = do
    numClients <- newMVar 0
    forever (do
        connection <- listenOn (Port 8080)
        forkIO (handleClient numClients connection) )

modifyMVar :: (a -> a) -> MVar a -> IO ()
modifyMVar f mVar = do
    v <- takeMVar mVar
    putMVar mVar $! f v

handleClient :: MVar Integer -> Handle -> IO ()
handleClient numClients connection =
    modifyMVar numClients (+1)
    getHttpMethod connection >= \case
        Head -> respond connection "I'm not RFC compliant weee\r\n\r\n"
        Get  -> ...
        Delete -> ...
        ...
    modifyMVar numClients (subtract 1)
    close connection
```

# Higher abstractions

- ▶ Channels: queues built with MVars
- ▶ Semaphores: more flexible locking
- ▶ Concurrent data structures

## Shortcomings of forkIO

- ▶ Silent exceptions
- ▶ Returning values from forks is awkward
- ▶ Locks are hard to use, no atomicity guarantees

STM

## Transactional memory

```
modifyMVar :: (a -> a) -> MVar a -> IO ()  
modifyMVar f mVar = do  
    value <- takeMVar mVar  
    putMVar mVar (f value)
```

What happens when another thread accesses the MVar between the take and put?

# MVar API

```
data IO a -- abstract  
instance Monad IO where ...
```

```
data MVar a -- abstract  
newMVar :: a -> IO (MVar a)  
takeMVar :: MVar a -> IO a  
putMVar :: MVar a -> a -> IO ()  
modifyMVar :: (a -> a) -> MVar a -> IO () -- Dangerous
```

```
block :: IO a -- doesn't exist, but might be handy sometimes  
whenBlockedThen :: IO a -> IO a -> IO a -- dito
```

```
throw :: Exception e => e -> IO a  
catch :: Exception e => IO a -> (e -> IO a) -> IO a
```

## TVar API

```
data STM a -- abstract
instance Monad STM where ...
```

```
data TVar a -- abstract
newTVar :: a -> STM (TVar a)
takeTVar :: TVar a -> STM a
putTVar :: TVar a -> a -> STM ()
modifyTVar :: (a -> a) -> TVar a -> STM ()
```

```
retry :: STM a
orElse :: STM a -> STM a -> STM a
```

```
throwSTM :: Exception e => e -> STM a
catchSTM :: Exception e => STM a -> (e -> STM a) -> STM a
```

```
atomically :: STM a -> IO a
```

# TVar API

```
data STM a -- abstract
instance Monad STM where ...
```

```
data TVar a -- abstract
newTVar :: a -> STM (TVar a)
takeTVar :: TVar a -> STM a
putTVar :: TVar a -> a -> STM ()
modifyTVar :: (a -> a) -> TVar a -> STM ()
```

```
retry :: STM a
orElse :: STM a -> STM a -> STM a
```

```
throwSTM :: Exception e => e -> STM a
catchSTM :: Exception e => STM a -> (e -> STM a) -> STM a
```

```
atomically :: STM a -> IO a
```

```
modifyTVar :: (a -> a) -> TVar a -> STM ()
modifyTVar f tVar = do
    value <- readTVar tVar
    writeTVar tVar (f value)
```



## How does it work?

- ▶ STM actions write to a log of changes
- ▶ atomically attempts to commit then log
- ▶ Changed circumstances lead to retrying the action
- ▶ Repeat until successful

# Should I use STM?

## Pros

- ▶ No deadlocks by design
- ▶ Guaranteed atomicity

## Cons

- ▶ Livelocks still possible
- ▶ Stampending herd
- ▶ No fairness guarantees
- ▶ Long transactions undesirable

## Speculative parallelism

“I think this is the answer  
so I’ll continue with that assumption  
and verify it in parallel”

... is a one-liner

```
spec :: Eq a => a -> (a -> b) -> a -> b  
spec g f a = let s = f g in s `par` if g == a then s else f a in spec
```

... is a one-liner

```
spec :: Eq a => a -> (a -> b) -> a -> b  
spec g f a = let s = f g in s `par` if g == a then s else f a in spec
```

- ▶ par is an old version of rpar from Control.Parallel
- ▶ Implemented by Edward Kmett
- ▶ ... after reading a 12-page paper attempting the same in C#
- ▶ ... written as a no-brainer one-liner in a Reddit comment

Add two minor modifications,

- ▶ Parametrize over the equality function
- ▶ No speculation with only one thread

*-- Copied verbatim from the `speculation` package*

```
specBy :: (a -> a -> Bool) -> a -> (a -> b) -> a -> b
specBy cmp guess f a
  | numCapabilities == 1 = f $! a
  | otherwise = speculation `par`
    if cmp guess a
      then speculation
      else f a
  where speculation = f guess

spec :: Eq a => a -> (a -> b) -> a -> b
spec = specBy (==)
```

## How does it work?

- ▶ Threads are heap objects
- ▶ Heap objects are tracked by the GC
- ▶ Obsolete threads are simply GC'd
- ▶ Fork as long as you have memory and CPU free



## Things not covered

- ▶ Exceptions
- ▶ Debugging
- ▶ `async` library
- ▶ FFI integration
- ▶ Distributed processes
- ▶ Data parallel computations
- ▶ GPGPU
- ▶ Other semiautomatic parallelization libs

Questions?