

Parallelism and concurrency in Haskell

David Luposchainsky

2017-12-07

Introduction

Concurrency vs. parallelism

Parallelism

- ▶ Goal: speedup
- ▶ Orthogonal to correctness
- ▶ Comparatively simple

Concurrency

- ▶ Affects logical structure of a program
- ▶ Heavily interleaved with correctness
- ▶ Often very complicated

Talk outline

1. **Parallelism**

Excellent case for laziness

2. **Concurrency with forks and mutable state**

Nice abstractions, language agnostic

3. **Transactional memory**

Unique to Haskell, envied by everyone else

Parallelism

“Go faster”

Checking primes

Problem: find primes in a list (e.g. for crypto)

```
isPrime :: Natural -> Bool
```

```
filterPrime :: [Natural] -> [Natural]
```

```
filterPrime = filter isPrime
```

```
primesBetween lo hi = filterPrime [lo..hi]
```

Checking primes

Problem: find primes in a list (e.g. for crypto)

```
isPrime :: Natural -> Bool
```

```
filterPrime :: [Natural] -> [Natural]
```

```
filterPrime = filter isPrime
```

```
primesBetween lo hi = filterPrime [lo..hi]
```

Unfortunately, this won't parallelize well.

`map` is simple enough to parallelize (why?). Can we use that fact?

map is simple enough to parallelize (why?). Can we use that fact?

```
isPrime :: Natural -> Bool
```

```
filterPrime :: [Natural] -> [Natural]
```

```
filterPrime = map snd . filter fst . map (\n -> (isPrime n, n))
```

```
    -- `isPrime` isolated in `map` argument
```

```
    -- (using the decorate-operation-undecorate idiom)
```

```
primesBetween lo hi = filterPrime [lo..hi]
```

What's in the list?

```
map (\n -> (isPrime n, n))  
    (take 1e5 [1e10..])
```

What's in the list?

```
map (\n -> (isPrime n, n))
    (take 1e5 [1e10..])
```

Thanks! Literally function applications of the lambda.

[illegible]

Forcing

- ▶ = Evaluation to a certain degree (e.g. *WHNF*)
- ▶ Primitive task that allows parallelism
- ▶ “I’ll probably need this, so maybe evaluate it”

Evaluation strategies

```
-- Control.Parallel.Strategies

-- A way to evaluate 'a's
type Strategy a = a -> Eval a
withStrategy :: Strategy a -> (a -> a)
```

Evaluation strategies

```
-- Control.Parallel.Strategies

-- A way to evaluate 'a's
type Strategy a = a -> Eval a
withStrategy :: Strategy a -> (a -> a)

-- Primitives
r0    :: Strategy a -- "NOOP"
rseq  :: Strategy a -- Sequential forcing
rpar  :: Strategy a -- Parallel forcing
```

Evaluation strategies

```
-- Control.Parallel.Strategies

-- A way to evaluate 'a's
type Strategy a = a -> Eval a
withStrategy :: Strategy a -> (a -> a)

-- Primitives
r0    :: Strategy a -- "NOOP"
rseq  :: Strategy a -- Sequential forcing
rpar  :: Strategy a -- Parallel forcing

-- Combining functions
evalList      :: Strategy a -> Strategy [a]
evalTraversable :: Traversable t => Strategy a -> Strategy (t a)
evalTuple2    :: Strategy a -> Strategy b -> Strategy (a,b)
```


Back to primes

```
filterPrime =  
  map snd . filter fst . map (\n -> (isPrime n, n))  
  --  
  -- List of thunks we'd like to force in parallel
```

Back to primes

```
filterPrime =  
    map snd . filter fst . map (\n -> (isPrime n, n))  
    --  
    -- List of thunks we'd like to force in parallel  
  
filterPrimeParallel =  
    map snd . filter fst . parallelize . map (\n -> (isPrime n, n))
```

Back to primes

```
filterPrime =
  map snd . filter fst . map (\n -> (isPrime n, n))
  --
  -- List of thunks we'd like to force in parallel

filterPrimeParallel =
  map snd . filter fst . parallelize . map (\n -> (isPrime n, n))

where
  parallelize :: [(a, b)] -> [(a, b)]
  parallelize = withStrategy tupleFstList

  tupleFstList :: Strategy [(a, b)]
  tupleFstList = parList parFst

  parFst :: Strategy (a,b)
  parFst = parTuple2 rseq r0
```

Back to primes

```
filterPrime =
  map snd . filter fst . map (\n -> (isPrime n, n))
  --
  -- List of thunks we'd like to force in parallel

filterPrimeParallel =
  map snd . filter fst . parallelize . map (\n -> (isPrime n, n))

where
  parallelize :: [(a, b)] -> [(a, b)]
  parallelize = withStrategy tupleFstList

  tupleFstList :: Strategy [(a, b)]
  tupleFstList = parList parFst

  parFst :: Strategy (a,b)
  parFst = parTuple2 rseq r0

filterPrimeParallel =
  map snd . filter fst . parallelize . map (\n -> (isPrime n, n))
  where
    parallelize = withStrategy (parList (parTuple2 rseq r0))
```

Demo

```
. ./run  
[1 of 1] Compiling Main ( Primes.hs, Primes.o )  
Linking Primes ...  
Running sequentially  
45.38 s  
Running in parallel  
16.37 s
```

Recap

`Control.Parallel` for easy parallelism:

1. Define strategy according to your needs
2. Apply strategy, done

Threads are cheap

- ▶ Couple of hundred (heap) bytes per thread
- ▶ Literally millions are possible
- ▶ Don't be afraid of using them!

```
>>> time ./Main.hs 1000000
```

```
Forking 1000000 threads
```

```
Start the counting cascade
```

```
Counter: 1000000
```

```
./Main.hs 1000000 4,57s user 0,50s system 100% cpu 5,032 total
```

Other parallelism libraries

- ▶ **monad-par**: Explicit dataflow networks for complicated (pure) computations
- ▶ **speculation**: the one-liner to solve all pure speculative parallelism needs

Concurrency

Basic primitives

```
-- Threads  
forkIO :: IO () -> IO ThreadId  
forkFinally :: IO a  
            -> (Either SomeException a -> IO ())  
            -> IO ThreadId
```

Basic primitives

```
-- Threads
forkIO :: IO () -> IO ThreadId
forkFinally :: IO a
              -> (Either SomeException a -> IO ())
              -> IO ThreadId

-- Futures/promises
async :: IO a -> IO (Async a)
wait :: Async a -> IO a
```

Basic primitives

```
-- Threads
forkIO :: IO () -> IO ThreadId
forkFinally :: IO a
              -> (Either SomeException a -> IO ())
              -> IO ThreadId

-- Futures/promises
async :: IO a -> IO (Async a)
wait :: Async a -> IO a

-- Inter-thread communication
data MVar a
newMVar  :: a -> IO (MVar a)
takeMVar :: MVar a -> IO a
putMVar  :: a -> MVar a -> IO ()
```

Example: network service

```
main :: IO ()
main = do
  numClients <- newMVar 0
  forever (do
    connection <- listenOn (Port 8080)
    forkFinally (\_ -> close connection)
               (handleClient numClients connection) )
```

Example: network service

```
main :: IO ()
main = do
    numClients <- newMVar 0
    forever (do
        connection <- listenOn (Port 8080)
        forkFinally (\_ -> close connection)
                    (handleClient numClients connection) )

handleClient :: MVar Integer -> Handle -> IO ()
handleClient numClients connection = do
    modifyMVar numClients (+1)
    getHttpMethod connection >= \case
        Head -> ...
        Get  -> ...
        Delete -> ...
    ...
    modifyMVar numClients (subtract 1)
```

Higher abstractions

- ▶ Channels: queues built with MVars
- ▶ Semaphores: more flexible locking
- ▶ Concurrent data structures

STM

Transactional memory

```
modifyMVar :: (a -> a) -> MVar a -> IO ()  
modifyMVar f mVar = do  
    value <- takeMVar mVar  
    putMVar mVar (f value)
```

Transactional memory

```
modifyMVar :: (a -> a) -> MVar a -> IO ()  
modifyMVar f mVar = do  
    value <- takeMVar mVar  
    putMVar mVar (f value)
```

What happens when another thread fills the MVar between the take and put?

MVar API

```
data IO a
instance Monad IO where ...

data MVar a
newMVar    :: a -> IO (MVar a)
takeMVar   :: MVar a -> IO a
putMVar    :: MVar a -> a -> IO ()
modifyMVar :: (a -> a) -> MVar a -> IO () -- Dangerous

block      :: IO a
whenBlockedThen :: IO a -> IO a -> IO a

throw :: Exception e => e -> IO a
catch :: Exception e => IO a -> (e -> IO a) -> IO a
```

TVar API

```
data STM a
instance Monad STM where ...
```

```
data TVar a
newTVar    :: a -> STM (TVar a)
takeTVar   :: TVar a -> STM a
putTVar    :: TVar a -> a -> STM ()
modifyTVar :: (a -> a) -> TVar a -> STM ()
```

```
retry  :: STM a
orElse :: STM a -> STM a -> STM a
```

```
throwSTM :: Exception e => e -> STM a
catchSTM :: Exception e => STM a -> (e -> STM a) -> STM a
```

```
atomically :: STM a -> IO a
```

TVar API

```
data STM a
instance Monad STM where ...

data TVar a
newTVar    :: a -> STM (TVar a)
takeTVar   :: TVar a -> STM a
putTVar    :: TVar a -> a -> STM ()
modifyTVar :: (a -> a) -> TVar a -> STM ()

retry  :: STM a
orElse :: STM a -> STM a -> STM a

throwSTM :: Exception e => e -> STM a
catchSTM :: Exception e => STM a -> (e -> STM a) -> STM a

atomically :: STM a -> IO a

modifyTVar :: (a -> a) -> TVar a -> STM ()
modifyTVar f tVar = do
    value <- readTVar tVar
    writeTVar tVar (f value)
```

How does it work?

- ▶ STM actions write to a log of changes
- ▶ atomically attempts to commit the log
- ▶ Changed circumstances lead to retrying the action
- ▶ Repeat until successful

Should I use STM?

Pros

- ▶ No deadlocks by design
- ▶ Guaranteed atomicity

Cons

- ▶ Livelocks still possible
- ▶ Stampending herd
- ▶ No fairness guarantees
- ▶ Long transactions undesirable

Speculative parallelism

“I think this is the answer
so I'll continue with that assumption
and verify it in parallel”

... is a one-liner

```
spec :: Eq a => a -> (a -> b) -> a -> b  
spec g f a = let s = f g in s `par` if g == a then s else f a in spec
```

... is a one-liner

```
spec :: Eq a => a -> (a -> b) -> a -> b  
spec g f a = let s = f g in s `par` if g == a then s else f a in spec
```

- ▶ par is an old version of rpar from Control.Parallel
- ▶ Implemented by Edward Kmett
- ▶ ...after reading a 12-page paper attempting the same in C#
- ▶ ...in a Reddit comment

Add two minor modifications,

- ▶ Parametrize over the equality function
- ▶ No speculation with only one thread

-- Copied verbatim from the `speculation` package

```
specBy :: (a -> a -> Bool) -> a -> (a -> b) -> a -> b
```

```
specBy cmp guess f a
```

```
  | numCapabilities == 1 = f $! a
```

```
  | otherwise = speculation `par`
```

```
    if cmp guess a
```

```
      then speculation
```

```
      else f a
```

```
  where speculation = f guess
```

```
spec :: Eq a => a -> (a -> b) -> a -> b
```

```
spec = specBy (==)
```

How does it work?

- ▶ Threads are heap objects
- ▶ Heap objects are tracked by the GC
- ▶ Obsolete threads are simply GC'd
- ▶ Fork as long as you have memory and CPU free

Things not covered

- ▶ Exceptions
- ▶ Debugging
- ▶ `async` library
- ▶ FFI integration
- ▶ Distributed processes
- ▶ Data parallel computations
- ▶ GPGPU
- ▶ Other semiautomatic parallelization libs

Questions?