

# Parallelism and concurrency in Haskell

David Luposchainsky

2017-12-07

## Introduction

# Concurrency vs. parallelism

## Parallelism

- ▶ Goal: speedup
- ▶ Orthogonal to correctness
- ▶ Comparatively simple

# Concurrency vs. parallelism

## Parallelism

- ▶ Goal: speedup
- ▶ Orthogonal to correctness
- ▶ Comparatively simple

## Concurrency

- ▶ Affects logical structure of a program
- ▶ Heavily interleaved with correctness
- ▶ Often very complicated

# Talk outline

1. **Parallelism**

Excellent case for laziness

2. **Concurrency with forks and mutable state**

Nice abstractions, language agnostic

3. **Transactional memory**

Unique to Haskell, envied by everyone else

Parallelism

“Go faster”

## Checking primes

Problem: find primes in a list (e.g. for crypto)

```
isPrime :: Natural -> Bool
```

```
filterPrime :: [Natural] -> [Natural]
```

```
filterPrime = filter isPrime
```

```
primesBetween lo hi = filterPrime [lo..hi]
```



## Checking primes

Problem: find primes in a list (e.g. for crypto)

```
isPrime :: Natural -> Bool
```

```
filterPrime :: [Natural] -> [Natural]
```

```
filterPrime = filter isPrime
```

```
primesBetween lo hi = filterPrime [lo..hi]
```

Unfortunately, this won't parallelize well.

`map` is simple enough to parallelize (why?). Can we use that fact?

map is simple enough to parallelize (why?). Can we use that fact?

```
isPrime :: Natural -> Bool
```

```
filterPrime :: [Natural] -> [Natural]
```

```
filterPrime = map snd . filter fst . map (\n -> (isPrime n, n))
```

```
    -- `isPrime` isolated in `map` argument
```

```
    -- (using the decorate-operation-undecorate idiom)
```

```
primesBetween lo hi = filterPrime [lo..hi]
```

What's in the list?

```
map (\n -> (isPrime n, n))  
    (take 1e5 [1e10..])
```

## What's in the list?

```
map (\n -> (isPrime n, n))
    (take 1e5 [1e10..])
```

Thanks! Literally function applications of the lambda.

[illegible]

## Forcing

- ▶ = Evaluation to a certain degree (e.g. *WHNF*)
- ▶ Primitive task that allows parallelism
- ▶ “I’ll probably need this, so maybe evaluate it”

## Evaluation strategies

```
-- Control.Parallel.Strategies

-- A way to evaluate 'a's
type Strategy a = a -> Eval a
withStrategy :: Strategy a -> (a -> a)
```

## Evaluation strategies

```
-- Control.Parallel.Strategies

-- A way to evaluate 'a's
type Strategy a = a -> Eval a
withStrategy :: Strategy a -> (a -> a)

-- Primitives
r0    :: Strategy a -- "NOOP"
rseq  :: Strategy a -- Sequential forcing
rpar  :: Strategy a -- Parallel forcing
```



# Evaluation strategies

```
-- Control.Parallel.Strategies

-- A way to evaluate 'a's
type Strategy a = a -> Eval a
withStrategy :: Strategy a -> (a -> a)

-- Primitives
r0    :: Strategy a -- "NOOP"
rseq  :: Strategy a -- Sequential forcing
rpar  :: Strategy a -- Parallel forcing

-- Combining functions
evalList      :: Strategy a -> Strategy [a]
evalTraversable :: Traversable t => Strategy a -> Strategy (t a)
evalTuple2    :: Strategy a -> Strategy b -> Strategy (a,b)
```

## Back to primes

```
filterPrime =  
  map snd . filter fst . map (\n -> (isPrime n, n))  
  --  
  -- List of thunks we'd like to force in parallel
```

## Back to primes

```
filterPrime =  
    map snd . filter fst . map (\n -> (isPrime n, n))  
    --  
    -- List of thunks we'd like to force in parallel  
  
filterPrimeParallel =  
    map snd . filter fst . parallelize . map (\n -> (isPrime n, n))
```

## Back to primes

```
filterPrime =
  map snd . filter fst . map (\n -> (isPrime n, n))
  --
  -- List of thunks we'd like to force in parallel

filterPrimeParallel =
  map snd . filter fst . parallelize . map (\n -> (isPrime n, n))

where
  parallelize :: [(a, b)] -> [(a, b)]
  parallelize = withStrategy tupleFstList

  tupleFstList :: Strategy [(a, b)]
  tupleFstList = parList parFst

  parFst :: Strategy (a,b)
  parFst = parTuple2 rseq r0
```

## Back to primes

```
filterPrime =
  map snd . filter fst . map (\n -> (isPrime n, n))
  --
  -- List of thunks we'd like to force in parallel

filterPrimeParallel =
  map snd . filter fst . parallelize . map (\n -> (isPrime n, n))

where
  parallelize :: [(a, b)] -> [(a, b)]
  parallelize = withStrategy tupleFstList

  tupleFstList :: Strategy [(a, b)]
  tupleFstList = parList parFst

  parFst :: Strategy (a,b)
  parFst = parTuple2 rseq r0

filterPrimeParallel =
  map snd . filter fst . parallelize . map (\n -> (isPrime n, n))
  where
    parallelize = withStrategy (parList (parTuple2 rseq r0))
```

## Demo

```
. ./run  
[1 of 1] Compiling Main ( Primes.hs, Primes.o )  
Linking Primes ...  
Running sequentially  
45.38 s  
Running in parallel  
16.37 s
```

## Recap

`Control.Parallel` for easy parallelism:

1. Define strategy according to your needs
2. Apply strategy, done

## Threads are cheap

- ▶ Couple of hundred (heap) bytes per thread
- ▶ Literally millions are possible
- ▶ Don't be afraid of using them!

```
>>> time ./Main.hs 1234567
```

```
Forking 1234567 threads
```

```
Start the counting cascade
```

```
Counter: 1234567
```

```
./Main.hs 1234567  5,38s user 0,66s system 100% cpu 6,001 total
```



## Other parallelism libraries

- ▶ **monad-par**: Explicit dataflow networks for complicated (pure or effectful) computations
- ▶ **speculation**: the one-liner to solve all pure speculative parallelism needs
- ▶ **Haxl**: See Simon Marlow's BTD 10 talk

## Concurrency

## Basic primitives

```
-- Threads  
forkIO :: IO () -> IO ThreadId  
forkFinally :: IO a  
            -> (Either SomeException a -> IO ())  
            -> IO ThreadId
```

## Basic primitives

```
-- Threads
forkIO :: IO () -> IO ThreadId
forkFinally :: IO a
              -> (Either SomeException a -> IO ())
              -> IO ThreadId

-- Futures/promises
async :: IO a -> IO (Async a)
wait :: Async a -> IO a
```

## Basic primitives

```
-- Threads
forkIO :: IO () -> IO ThreadId
forkFinally :: IO a
              -> (Either SomeException a -> IO ())
              -> IO ThreadId

-- Futures/promises
async :: IO a -> IO (Async a)
wait :: Async a -> IO a

-- Locking
data MVar a
newMVar  :: a -> IO (MVar a)
takeMVar :: MVar a -> IO a
putMVar  :: a -> MVar a -> IO ()
```

## Example: webserver

```
main :: IO ()
main = do
    numClients <- newMVar 0
    forever (acceptClient numClients)

acceptClient :: MVar Int -> IO ()
acceptClient numClients = do
    connection <- listenOn (Port 8080)
    forkFinally (\_ -> close connection)
                (handleClient numClients connection)
```

## Example: webserver

```
main :: IO ()
main = do
    numClients <- newMVar 0
    forever (acceptClient numClients)

acceptClient :: MVar Int -> IO ()
acceptClient numClients = do
    connection <- listenOn (Port 8080)
    forkFinally (\_ -> close connection)
                (handleClient numClients connection)

handleClient :: MVar Int -> Handle -> IO ()
handleClient numClients connection = do
    modifyMVar numClients (+1)
    getHttpMethod connection >>= \case
        Head    -> ...
        Get      -> ...
        Delete   -> ...
        ...
    modifyMVar numClients (subtract 1)
```

## Example: service chaining

```
main = do
  let customerId = 123
  addressA <- async (queryAddress customerId)
  inventoryA <- async (queryInventory customerId)
```



## Example: service chaining

```
main = do
  let customerId = 123
  addressA <- async (queryAddress customerId)
  inventoryA <- async (queryInventory customerId)

  address <- wait addressA
  capabilitiesA <- async (queryCapabilities customerId address)
```

## Example: service chaining

```
main = do
  let customerId = 123
  addressA <- async (queryAddress customerId)
  inventoryA <- async (queryInventory customerId)

  address <- wait addressA
  capabilitiesA <- async (queryCapabilities customerId address)
  capabilities <- wait capabilities

  -- more stuff here

  inventory <- wait inventoryA
  doStuff address inventory address
```

## Higher abstractions

- ▶ Channels: queues built with MVars
- ▶ Semaphores: more flexible locking
- ▶ Concurrent data structures

STM

## Transactional memory

```
modifyMVar :: (a -> a) -> MVar a -> IO ()  
modifyMVar f mVar = do  
    value <- takeMVar mVar  
    putMVar mVar (f value)
```

## Transactional memory

```
modifyMVar :: (a -> a) -> MVar a -> IO ()  
modifyMVar f mVar = do  
    value <- takeMVar mVar  
    putMVar mVar (f value)
```

What happens when another thread fills the MVar between the take and put?

## Standard API

```
data IO a
instance Monad IO where ...
```

```
data MVar a
newMVar    :: a -> IO (MVar a)
takeMVar   :: MVar a -> IO a
putMVar    :: MVar a -> a -> IO ()
modifyMVar :: (a -> a) -> MVar a -> IO () -- Dangerous
```

# Transactional API

```
data STM a
instance Monad STM where ...

data TVar a
newTVar    :: a -> STM (TVar a)
takeTVar   :: TVar a -> STM a
putTVar    :: TVar a -> a -> STM ()
modifyTVar :: (a -> a) -> TVar a -> STM ()

retry  :: STM a
orElse :: STM a -> STM a -> STM a

atomically :: STM a -> IO a
```



# Transactional API

```
data STM a
instance Monad STM where ...

data TVar a
newTVar      :: a -> STM (TVar a)
takeTVar     :: TVar a -> STM a
putTVar      :: TVar a -> a -> STM ()
modifyTVar   :: (a -> a) -> TVar a -> STM ()

retry  :: STM a
orElse :: STM a -> STM a -> STM a

atomically :: STM a -> IO a

modifyTVar :: (a -> a) -> TVar a -> STM ()
modifyTVar f tVar = do
    value <- readTVar tVar
    writeTVar tVar (f value)
```

## How does it work?

- ▶ STM actions write to a log of changes
- ▶ atomically attempts to commit the log
- ▶ Changed circumstances lead to retrying the action
- ▶ Repeat until successful

# Should I use STM?

## Pros

- ▶ No deadlocks by design
- ▶ Guaranteed atomicity

## Cons

- ▶ Livelocks still possible
- ▶ Stampending herd
- ▶ No fairness guarantees
- ▶ Long transactions undesirable

Speculative parallelism

“I think this is the answer  
so I'll continue with that assumption  
and verify it in parallel”

...is a one-liner

```
spec :: Eq a => a -> (a -> b) -> a -> b  
spec g f a = let s = f g in s `par` if g == a then s else f a in spec
```

...is a one-liner

```
spec :: Eq a => a -> (a -> b) -> a -> b  
spec g f a = let s = f g in s `par` if g == a then s else f a in spec
```

- ▶ par is an old version of rpar from Control.Parallel
- ▶ Implemented by Edward Kmett
- ▶ ...after reading a 12-page paper attempting the same in C#
- ▶ ...in a Reddit comment

Add two minor modifications,

- ▶ Parametrize over the equality function
- ▶ No speculation with only one thread

*-- Copied verbatim from the `speculation` package*

```
specBy :: (a -> a -> Bool) -> a -> (a -> b) -> a -> b
```

```
specBy cmp guess f a
  | numCapabilities == 1 = f $! a
  | otherwise = speculation `par`
    if cmp guess a
      then speculation
      else f a
  where speculation = f guess
```

```
spec :: Eq a => a -> (a -> b) -> a -> b
```

```
spec = specBy (==)
```



## How does it work?

- ▶ Threads are heap objects
- ▶ Heap objects are tracked by the GC
- ▶ Obsolete threads are simply GC'd
- ▶ Fork as long as you have memory and CPU free

## Things not covered

- ▶ Exceptions
- ▶ Debugging
- ▶ FFI integration
- ▶ Distributed processes
- ▶ Data parallel computations
- ▶ GPGPU
- ▶ Other semiautomatic parallelization libs

Questions?